

CHAPTER 3 km reanalysis — known properties of the archive

Third of the three documents that ship with this archive. `CHAPTER_VARIABLES.md` says what is in the files; `MISSING_VARIABLES.md` says what is not and why. This one answers the question neither of those asks: **what is in the files that will surprise you**. Everything below is a property of the data as delivered. None of it is a defect awaiting repair: the defects the audit found were repaired before the archive was produced, and they are not listed here. What is listed is the archive behaving correctly in ways a user comparing it against ERA5, against observations, or against their own expectations will otherwise discover alone — and, having discovered one, will reasonably stop trusting the rest. Every number here was measured, not estimated. They come from a variable-by-variable audit run in September 2026 over a **34-timestep sample** spanning both sampled years and all four seasons (the sample is listed in `docs/audit/sample.md`); where a comparison against ERA5 is quoted, our 3 km field was upscaled to the ERA5 grid first. Each entry links to the issue that holds the measurement and the method.

How to read this document

Section 0 is the short list: the things that change what you do, in one line each. The rest states each property in full, with its number and what to do about it. Two kinds of entry are mixed together deliberately, because a user does not care which is which:

- **properties of the run** — the model was configured a certain way and the data reflect it. Nothing can be done after the fact; the run would have to be repeated.
- **properties of the conversion** — a deliberate convention, chosen and documented. These could be changed, and the entry says what was chosen instead and why.

Which archive this describes

The GRIB2 archive, `grib_v3`, **246 messages per hourly file**, as produced by the converter from 2026-09-25 onward. Files written before that date carry six defects that were repaired on that date — relative humidity, specific humidity, the column integrals, snowfall, the water masks and convective inhibition. They are listed in the repository's audit record (`docs/audit/re-verification.md`) and are **not** properties of this archive. If you hold files from before the rebuild, replace them rather than reconciling them.

0. The short list

#	Property	The number
1.1	Surface shortwave is high under clear sky: the run carried no aerosol	+8.5 % clear-sky, +13.6 % all-sky vs ERA5
1.2	Total cloud cover is low, high cloud lowest — it is CLDFRA, not the overlap	0.78 and 0.61 of ERA5
1.3	Land skin temperature is cold, and the archive cannot say why	−1.64 K mean, −3.24 K at midday
2.1	<code>fsr</code> is WRF's local roughness, ERA5's is effective	0.32-0.48 of ERA5 over land
2.2	<code>zust</code> over land responds to resolved terrain	1.6-2.1× ERA5 over land, 1.1× over sea

#	Property	The number
3.1	r and 2r saturate differently, and each is right	mixed phase vs liquid water
3.2	r above 100 % is normal and is not clipped	up to 139.7 % at 300 hPa
3.3	Humidity rebuilt from 2t and 2d overshoots at saturation	peaks at 100.09 %
4.1	Below ground the pressure levels are clamped , not extrapolated — ERA5 extrapolates	affects 47.6 % of the domain at 1000 hPa
4.2	z is exactly $9.80665 \times z_{\text{WRF}}$; WRF integrated with 9.81	0.034 % off hydrostatic balance
5.1	10fg is the resolved 10 m wind, not a gust parameterisation	0.625 of ERA5's gust
5.2	Its declared one-hour window is right 88 % of the time	excess 0.499-0.784 m/s on the bad hours
6.1	Three runoff variables carry one field's information	ssro non-zero at $\leq$ 19 points
6.2	tcw sums six species where ECMWF defines five	graupel residual to 26 kg/m²
7.1	One masking rule: a point is missing where the ERA5 parameter is undefined	bitmaps, never zeros
7.2	mucin is masked , not zero-filled, where the scheme declined	missing on 71-96 % of the domain
8.1	generatingProcessIdentifier = 128 marks every message and discriminates nothing	246 of 246
8.2	al is a two-season table value by category, not a continuous field	8-11 distinct values domain-wide
8.3	slor is the slope of the resolved 3 km terrain, not of sub-grid orography	eccodes shows the ECMWF name

1. Radiation, cloud and the surface energy budget

These three belong together. Read any one of them alone and you will predict the other two wrongly.

1.1 Surface shortwave is 8.5 % high under clear sky, because the run carried no aerosol

RRTMG ran with a CAM ozone climatology and **no aerosol at all** (aer_opt = 0). ERA5 carries the Tegen aerosol climatology. Over this domain — Europe, the Mediterranean and the Saharan margin — an aerosol optical depth of 0.1 to 0.2 attenuates the direct beam by about the amount measured. Ratio of means against ERA5 over the sample:

field	ours / ERA5
ssrdc surface solar down, clear sky	1.085
ssrd surface solar down, all sky	1.136
ssrc surface net solar, clear sky	1.116
ssr surface net solar, all sky	1.178
tsrc top net solar, clear sky	1.031
strd, strdc downward thermal	0.989, 0.996
ttr, ttrc outgoing thermal	1.009, 0.990

The decomposition is the useful part: the clear-sky figure isolates the atmosphere with no cloud involved, and the thermal fields — untouched by aerosol — agree with ERA5 to within 1.1 %, which is the control that makes the attribution credible rather than a guess. **What to do.** If you use this archive for solar resource, for an energy application, or as training data against pyranometer observations, expect a systematic high bias in surface shortwave of about this size. It is a property of the run and cannot be removed after the fact. Measured in #45.

1.2 Total cloud cover runs 22 % below ERA5, and high cloud 39 % below

Sample means, upscaled to the ERA5 grid:

	ours	ERA5	ratio
tcc	0.381	0.487	0.783
hcc	0.169	0.276	0.610
mcc	0.161	0.205	0.786
lcc	0.231	0.268	0.862
tcw	0.0185	0.0166	1.109

The deficit holds at **every one of the 34 timesteps**, both years, every month. **It is not the overlap assumption, and that was checked first.** The converter uses maximum-random overlap where ECMWF uses exponential-random, which would yield more cover — the right direction. But random overlap is the *most* total cover any assumption can extract from a given profile, and computed point by point on our own CLDFRA the entire span of possible assumptions is only **2 to 4 % wide**, because CLDFRA is strongly bimodal. ERA5 is 28 % above us. The overlap accounts for at most a seventh of the gap. The gap is in CLDFRA itself — WRF's Xu-Randall diagnostic under ICLoud = 1. The sharpest statement of it: **we carry 11 % more cloud ice than ERA5 and diagnose 39 % less high cloud.** The condensate is there; the fraction the scheme assigns to it is not. **One caveat, stated rather than buried.** Part of the gap is resolution. Our tcc is overlapped at 3 km and then averaged to 31 km; ERA5's is overlapped on an already-smooth 31 km profile, which yields somewhat more cover. That effect is bounded by the same 2–4 % span, so it cannot account for 22 % — but it is real, and the figure should not be quoted as if it were all model error. Measured in #48.

1.3 Land skin temperature runs 1.6 K below ERA5, 3.2 K at midday — and this contradicts the two above

skt against ERA5, mean over the whole sample:

region	bias	correlation
sea	+0.01 K	0.95–0.98
land	–1.64 K , reaching –3.24 K at midday	0.95–0.98

At 00Z the land bias is -0.5 to -1.2 K; at 12Z it is -0.8 to -3.2 K. **The sea number is the control.** Over open water the model's skin temperature is the SST, and both reanalyses see nearly the same one. A bias of $+0.01$ K there says the derivation, the emissivity and the encoding are not the cause of anything (see 2.3): the land bias is the run. **And it pulls against §1.1 and §1.2.** More shortwave reaching the ground and less cloud overhead both put energy *into* the land surface, and the surface comes out colder — most so at the hour when the excess solar is largest. That is not a contradiction in the measurements; it is an unexplained partition, and reconciling it needs the surface energy budget. **HFX, LH and the ground heat flux are precisely the three fields this archive does not have** (see MISSING_VARIABLES.md), so it cannot be settled from the published data. Candidates that cannot be separated: the land-surface scheme partitioning more of the available energy into latent heat, a different soil-moisture climate, the resolved 3 km orography raising the mean elevation of an upscaled land cell, or a genuine cold bias. **What to do.** If you use skt over land for evaluation, bias correction, or ML training against satellite land surface temperature, apply the number. Do not predict a warm land surface from the radiation findings — the archive does the opposite. Measured in #57.

2. Surface exchange

2.1 f_{sr} is a local roughness where ERA5's is effective

	ours	ERA5	ratio of medians
f _{sr} , land	0.188–0.217 m	0.375–0.394 m	0.32–0.48
f _{sr} , sea	0.0089–0.0091 m	0.0103–0.0105 m	1.5–1.9

ERA5's f_{sr} is an **effective** roughness that carries the orographic drag the IFS parameterises at 31 km. WRF's ZNT, which is what this field is, is the **local** surface roughness of the land-use class — the orography is resolved at 3 km instead of parameterised. Over sea the two agree, because both are Charnock. ZNT is a land-use lookup with a snow override: one value per category — 0.70 m for the forest classes, 0.15 for savannas and grassland, 0.065 for barren, exactly 1.00 m on the 19 867 urban points — with a floor of 0.011 m where snow covers the surface. Nothing is capped or mis-scaled. **What to do.** Do not take f_{sr} from this archive and f_{sr} from ERA5 as the same quantity: over land they differ by a factor of two to three. A drag coefficient computed from ours and compared against one computed from ERA5's will show a discrepancy that is not physical. Measured in #56.

2.2 z_{ust} runs 1.6 to 2.1 times ERA5's over land, and agrees over sea

	ours	ERA5	ratio
z _{ust} , land	0.44–0.52 m/s	0.23–0.28 m/s	1.62–2.06
z _{ust} , sea	0.21–0.34 m/s	0.19–0.31 m/s	1.08–1.12 , corr 0.85–0.90

This is the other half of §2.1's story, and the split is what identifies it: our roughness is a third of ERA5's while our friction velocity is twice it, and those cannot both be a scaling error on the same z₀. At 3 km the terrain-driven wind is resolved rather than smoothed, and the friction velocity responds to it. The excess is largest at midday — 1.06–1.13 at 00Z against 1.40–1.75 at 12Z, in every region. The field is internally consistent with the published stresses: $|\tau| = \rho u^*$ holds to $1.8e-6$ N/m² and $z_{ust} = \sqrt{|\tau|/\rho}$ to $2.7e-4$ m/s, at every point of every timestep checked. Measured in #56.

2.3 skt is derived, its accuracy is known, and it is seasonal

TSK was not written out by the run, so skt is inverted from the upward longwave flux: $LWUPB = \epsilon \sigma T^4 + (1-\epsilon) GLW$, with the per-category emissivity **measured from the run itself** rather than looked up in a table. The derivation is verified; what is seasonal is how well it reproduces the model's own state:

	figure
water gate — skt against the model's own SST	0.0005 K RMSE , absmax 0.002 K, ~910 000 points, every timestep
land points reproduced to 1e-4	99.98 % in July ... 93.05 % in January
land skt RMSE from the emissivity table	0.0012 K in July ... 0.115-0.133 K in January and February , p99 0.85 K, absmax 2.8 K

The whole of the seasonal spread is the snow rule: emissivity switches to a flat 0.98 under snow, and the switch is applied slightly earlier than the model's own transition completes. Over water, where the answer is known independently, the derivation is exact. **What to do.** Treat skt over water as the model's SST to a thousandth of a kelvin. Over land in winter, allow a tenth of a kelvin of derivation error on top of the physical bias of §1.3 — the two are unrelated and the derivation error is thirty times smaller. Measured in #55.

3. Humidity — read this before using r or 2r

3.1 r and 2r saturate against different phases, and each is right for what it is

This is the single most likely misreading of the archive, so it gets the plainest sentence available:

r (paramId 157, pressure levels) is relative humidity over the mixed phase. 2r (paramId 260242, 2 m) is relative humidity over liquid water at every temperature. They are not the same quantity and will not agree at cold points.

The reason is that each follows its own definition rather than a single house convention:

- r is ECMWF's paramId 157, which saturates over ice below 250.16 K, over water above 273.16 K, and quadratically between. The archive computes it from IFS Cy41r2's own equations (7.5, 7.6, 7.89, 7.90) with every constant taken from that cycle — the one ERA5 was produced with — so it is directly comparable with ERA5's r. Over ice the mixed-phase and over-water definitions differ by up to a factor **1.8**.
- 2r follows the **WMO screen convention**, which saturates over liquid water at all temperatures. ERA5 publishes no 2 m relative humidity, so there is no ERA5 field to follow; the choice is the meteorological standard for a screen-level observation, which is what a 2 m humidity is compared against.

What to do. Do not build a profile that puts 2r at the bottom of r. Do not compute a vertical humidity gradient across the two. If you need them on one convention, recompute the one you need from q, t and sp — all three are published. Decided and measured in #40.

3.2 r above 100 % is normal, not a defect

There is no clip. Supersaturation is physical, the model produces it, and ERA5 publishes it too — so the archive publishes it. Measured on a single file: **139.7 % at 300 hPa** in the ice regime, and **34 785 points above 100 %** at 400 hPa. Below ground, values reach further still — up to about 163 % — and that is the clamping of §4.1 rather than the humidity: the lowest model level's moisture, from roughly 25 m above the terrain, read at a nominal pressure of 1000 hPa. See §4.1 before treating it as a moisture error. If your ingest pipeline rejects relative humidity above 100, it will reject valid data from this archive. Measured in #40 and #44.

3.3 Humidity reconstructed from 2t and 2d peaks at 100.09 %

At saturated points 2d exceeds 2t by a fixed **0.0133 K**, on 6 000 to 64 000 points per timestep. Every violating point is at relative humidity 100 — it is the dewpoint inversion overshooting by a constant at exactly saturation, and the constant is the same number at every timestep in every season, which is what identifies it as arithmetic rather than noise. It was left in place deliberately. Clipping 2d to 2t would remove the visible crossings while leaving the same bias at every unsaturated point, put a kink in the field at RH 100, and break the exact consistency between 2d and Q2 that the audit verified at $3.6\text{e-}4$ to $1.1\text{e-}2$ in vapour pressure. That trades a constraint for the appearance of one. **What to do.** Expect a reconstruction that peaks a tenth of a per cent above saturation. If you need a hard bound, clip your own reconstruction rather than reading it as an inconsistency in the archive. Measured in #42.

4. Vertical structure

4.1 Below ground the pressure levels are clamped to the lowest model level — ERA5 extrapolates

Where a pressure level lies below the terrain, this archive publishes **the lowest model level's value, verbatim**. It is not an extrapolation and not a fill: t matches the lowest model level to **exactly 0.0** at 100 % of below-ground points, q to $1\text{e-}9$, z to $2.5\text{e-}4$ m. There is a good half and a bad half, and both matter. **The good half.** The number is a real model state from about 24 to 27 m above the ground, not an invention. A user reading 1000 hPa over the Alps is reading near-surface air, which is a more defensible thing to publish than a synthetic lapse-rate extrapolation. It is also why the archive is coherent across the surface join: $2t - t(1000 \text{ hPa})$ below ground is -0.9 K at midnight and $+1.4 \text{ K}$ at noon in July, which is a nocturnal inversion and a superadiabatic daytime surface layer over a 22 m interval. **The bad half.** z then contradicts sp . Where the surface pressure says the 1000 hPa surface is underground, the published z at 1000 hPa puts it a median of **24 to 27 m above** the terrain. Both messages are in the same file. Anyone deriving a height from z , or a thickness between levels, meets that self-contradiction on **47.6 % of the domain at 1000 hPa** and 9 % at 925 hPa. **And ERA5 does not do this.** ERA5 extrapolates below ground with a lapse rate, so the two archives differ in *convention* exactly where 1000 hPa is underground — which is most of the land. This difference, not any defect, is the whole of the 1000 hPa discrepancy against ERA5: a t bias of -1.71 K at 1000 hPa against -0.05 K at 850 hPa, and a z correlation of 0.108. **What to do.** Mask the pressure levels against sp yourself if you need a clean comparison against ERA5 or a thickness computation. Every field needed to do it is in the same file. Measured in #44.

4.2 z is exactly $9.80665 \times \text{WRF's height}$, and the run integrated with 9.81

Both halves of this are true and a user will find one of them alone:

- **By construction**, the published geopotential is standard gravity times WRF's own geopotential height. Divide z by 9.80665 and you recover the model's height in metres, exactly — at the surface and on every pressure level, with no residual.
- **WRF integrated the hydrostatic equation with 9.81.** So the same archive sits **0.034 %** off hydrostatic balance against its own t and p . That is 1 part in 2 900: about 2 m at 500 hPa, 7 m at 50 hPa.

Neither number can be removed without breaking the other. The convention chosen is the one that makes z a correct ECMWF geopotential and makes the surface and pressure-level geopotentials agree with each other exactly — which they do, verified point by point over 1 349 910 land points. Measured in #39.

5. Wind

5.1 10fg is the resolved 10 m wind, not a gust parameterisation

comparison	ratio	corr
ours vs ERA5 10fg (gust scheme, hourly max)	0.625	0.813
ours vs ERA5 i10fg (gust scheme, instantaneous)	0.644	0.820
ours vs our own instantaneous 10 m wind, same file	1.004-1.032	—

The last row settles what the field is. 10fg exceeds the instantaneous 10 m wind speed of the same file by **0.5 to 3.2 %** in the mean, and is below it at zero points out of 2 220 273 — the consistency check an hourly maximum must pass, and it passes. So the “maximum” adds almost nothing: published under paramId 49, this field is **the 10 m wind speed with a two-per-cent premium**, where ERA5’s is a parameterised gust roughly 1.6 times the mean wind. WRF has no gust scheme here; this is the best proxy the run can give. The offset is a definition, not weather: 0.593 to 0.673 across all 34 sample timesteps, both years, every month, every hour. **What to do.** A user reading 10fg as a 10 m wind gust — which is what paramId 49 means everywhere else — **under-predicts by 38 %**. The correlation of 0.81 means it ranks windy places correctly; the ratio of 0.625 means it will not give you a gust threshold. For wind risk, insurance, infrastructure or training against station gust observations, treat it as a different quantity rather than a biased one. Measured in #50.

5.2 And its declared one-hour window is right only 88 % of the time

10fg is encoded as a one-hour maximum — stepType = max, stepRange = 0-1, typeOfStatisticalProcessing = 2, reference time H-1 — and that encoding was verified correct. **The field is not always a one-hour maximum.** WRF resets WSPD10MAX at the history write, and measured over six whole days, 138 hourly transitions, **the reset was missed at 17 of them (12.3 %)**. Those hours carry a maximum over two hours or more under a message that declares one. How much it costs, as mean excess over the instantaneous 10 m wind of the same file:

	mean excess
hours that reset	0.068-0.311 m/s
hours that did not	0.499-0.784 m/s

Three to eight times larger — a 10 to 17 % overstatement on a domain mean of about 4.5 m/s, at hours whose message says one hour. **Worse, it moves those hours closer to ERA5’s gust for the wrong reason**, so the error is not detectable from the value. It was left as 0-1 and stated here rather than repaired: detecting the carry-over requires the previous hour’s wrfout, which is not on disk when a file is converted, and no intra-file detector exists — all seven of WRF’s diagnostic maxima reset on the same alarm, so when it is missed they carry over together and none can witness for another. The hourly misses correlate with adaptive-timestep variability, but the mechanism was not established. **What to do.** Do not treat the 10fg time window as exact when a single hour matters. Over a season or a distribution the effect is the 12 % contamination above, which is quantified; for a specific hour there is no flag in the file. Measured in #49.

6. Water cycle

6.1 Three runoff variables carry one field's information

timestep	sro mean (m)	ssro mean (m)	ssro non-zero points	ro mean (m)
2024-02-15T12	3.093e-5	0	0 of 2 220 273	3.093e-5
2024-07-15T12	2.347e-5	2.7e-8	19	2.350e-5
2024-07-15T18	1.064e-4	3.9e-8	19	1.065e-4
2024-08-15T12	1.453e-5	0	0	1.453e-5

ssro is identically zero on five of the eighteen non-00Z sample timesteps, and where it is not, its domain mean is four to eight orders of magnitude below sro. **ro equals sro to four significant figures at every timestep.** The identity $ro = sro + ssro$ holds to $3.7e-9$ m; there is simply nothing in the second term. The cause is the time reference, not the conversion. The model's sub-surface drainage counter is alive — non-zero at 462 to 8 363 points with maxima of 11 to 106 mm — but it barely moves *within a day*: on 2024-07-15 the same 1 647 points carry drainage at 00, 06, 12 and 18Z and only 19 of them grow. The published ssro is the accumulation since 00Z, so the daily field is almost always empty. Against ERA5 the partition is the opposite way round: sro at 1.0–4.6 of ERA5's, ssro at 0.000–0.004, ro at 0.19–0.85. **What to do.** Do not close a water budget on the assumption that this archive has a drainage term — it does not, at daily scale. ro and sro are, knowingly, duplicate messages. We are not short of surface runoff; we are missing the other half of the partition, and ERA5's dominant term is the one we do not carry. Measured in #51.

6.2 tcw sums six species where ECMWF defines five

ECMWF defines paramId 136 tcw as the sum of water vapour, liquid water, cloud ice, rain and snow — **five species**. This archive builds it from **six**, adding graupel, and there is no graupel column published separately because ERA5 has no parameter for one. The residual, tcw minus the five published components, against the independently computed graupel column:

timestep	residual (mean)	graupel column (mean)	residual max
2024-01-15T12	0.003525	0.003396	5.69 kg/m²
2024-04-15T12	0.010819	0.010863	13.89 kg/m²
2024-07-15T12	0.009762	0.010219	26.46 kg/m²
2024-10-15T12	0.007046	0.007267	16.73 kg/m²

In the domain mean it is 0.4 % of the condensate and invisible. **Inside a July convective cell it reaches 26 kg/m² — larger than the whole of the rest of the column's condensate.** Keeping graupel inside tcw was the deliberate choice: dropping it would make tcw match ECMWF's definition exactly and the closure identity hold, at the cost of the graupel mass leaving the archive altogether. **What to do.** If you reconstruct tcw from tcwv + tclw + tciw + tcrw + tcsw, you will not get tcw, and the difference is graupel. In convective conditions it is not small. Measured in #47.

7. Masks and missing values

7.1 One masking rule for the whole archive

A point is masked where the ERA5 parameter is not defined there — not where the model has nothing to say, and not where the number looks wrong. Masked points are a GRIB

bitmap, never zeros, so a consumer can always tell "not defined here" from "zero here". Applied:

field	masked where
sst, ci	not sea
cl (zero) and dl (missing)	sea
swvl1-4, stll1-4	not permanent land
tv1	shrubland dominates — MODIS carries no shrub phenology and guessing was refused
slt	not permanent land
mucin	see §7.2

Two consequences are worth stating on their own. **The sea/inland-water distinction is not the land-use lake class.** WPS's MODIS class 21 calls the Black Sea, the Sea of Azov and part of the Baltic a lake — together 89 % of the cells this domain marks fully lake. Masking on it would have deleted sst from the Black Sea. The discriminator is connected-component **size**: the water components here are 844 431 cells (Atlantic + Mediterranean + Baltic + North Sea), 49 841 (Black Sea + Azov), 8 583 (Red Sea), then 1 225 (Vänern, the largest genuine inland lake) and smaller. The threshold sits in that gap with a factor of seven of slack. **The soil columns are masked off sea ice, and that discards something.** WRF reclassifies a water cell carrying sea ice as land, gives it soil type 16 and fills the column with water — swvl reaches exactly 1.000 against that soil type's porosity of 0.435. Those columns are the model's own state, but they are not soil, and ERA5 has no soil there because its land-sea mask is permanent. They are masked. Nothing is lost that cannot be recovered: ci is published, so the ice itself is in the file. Note the related fact that **lsm and al are not constant in time** in this archive — they follow the model hour by hour, because that is what it integrated, and WRF moves the land mask where sea ice appears. They are written in every message and must not be treated as one-per-archive statics. Decided and measured in #34, #36 and #38.

7.2 mucin is masked where the scheme declined, on most of the domain

The CAPE routine flags convective inhibition as missing wherever CAPE is below 100 J/kg. The archive writes a bitmap there rather than a zero, because zero would read as "nothing is stopping convection" at precisely the columns where the most is — a column with no CAPE is usually a column with a great deal of inhibition, and that is often *why* there is no CAPE. How much of the field is absent:

timestep	mucin present on	masked
2024-01-15T12	3.8 %	96.2 %
2024-02-15T12	5.7 %	94.3 %
2024-03-15T12	4.8 %	95.2 %
2024-07-15T12	28.8 %	71.2 %
2024-10-15T12	15.2 %	84.8 %
2025-02-15T12	5.5 %	94.5 %

On a winter day, 96 % of the mucin field is missing. Where we are silent, ERA5 carries an inhibition averaging 25 to 265 J/kg. And even where ours is present, it is 0.03 to 0.24 of ERA5's — because ours exists only in columns with CAPE ≥ 100 , which are the least inhibited ones, while ERA5 reports the field more widely. **What to do.** Expect a sparse field, most sparse in winter. Do not fill the gaps with zero: that is the state the archive deliberately does not publish, and a model trained on it learns that low-CAPE situations have no convective inhibition, which is the inverse of the physics. mucape is dense and its zeros are real — no CAPE genuinely is no available energy. Decided and measured in #54.

8. Conventions that do not mean what they look like

8.1 `generatingProcessIdentifier = 128` marks every message in the file

Within `grib_v3` this key **discriminates nothing**: all 246 messages carry 128, because that is ec-codes' own GRIB2 sample default and the converter sets it explicitly only for the accumulations. A consumer filtering on `generatingProcessIdentifier == 128` to find the 00Z-referred accumulations selects the entire file, including every instantaneous field. Nothing in the data is wrong — the accumulations *are* referred to 00 UTC and 128 does say so. It simply says it of everything else too. The value's only real meaning is against the older GRIB1 archive, where 127 marks the run-init-referred messages. **What to do.** Test `stepType / typeOfStatisticalProcessing` and `stepRange`, which state it unambiguously and were both verified. See §3 of `CHAPTER_VARIABLES.md` for the time convention itself. Measured in #53.

8.2 `al` is a two-season table value by land-use category, not a continuous field

The run used no monthly satellite albedo (`usemonalb = .false.`), so the background albedo comes from the land-use table's WINTER and SUMMER columns by dominant category. Consequences:

- the field takes only **8 to 11 distinct values** over the whole domain;
- it changes at ****1 171 937** land points — 89.9 % of the land mask — at the winter/summer boundary**, and every transition is a table row with no residue (0.20→0.17 croplands on 497 649 points, 0.23→0.25 barren on 369 219, and so on);
- SUMMER applies April to October and WINTER November to March, and the field does not move within a day;
- off land it is 0.08 on water and **0.65 on sea ice**, the model's default ice albedo.

What to do. Do not read `al` as an observed or evolving albedo, and do not difference it across the seasonal boundary expecting a physical signal. Note also that it moves *twice* for another reason: at sea-ice points, following the land mask of §7.1. Measured in #37.

8.3 `s1or` is the slope of the resolved terrain, and `sdor` is ERA5's quantity at 3 km

Two orography parameters, differing from ERA5 in two different ways:

- `s1or` is published as the gradient magnitude of the **resolved** 3 km terrain, while `paramId 163` is defined as the slope of the **sub-grid** orography. ec-codes will show you the parameter's own name, *Slope of sub-gridscale orography*, and nothing in the message says otherwise — which is why it is stated here.
- `sdor` is the ERA5 quantity, evaluated on a 3 km cell instead of a 31 km one. The measured ratio of medians against ERA5 is 0.244, which is $(3/31)^{0.605}$ — a resolution rescaling, not a different definition.

Measured in #37.

8.4 "No ERA5 parameter fits" usually means "ERA5 does not publish it"

Several model fields are absent from this archive under a reason recorded as *no ERA5 parameter fits*. For most of them the accurate statement is different, and the distinction matters to anyone deciding whether to ask for the field:

field	the accurate reason
SNOALB	ECMWF defines 260161 <code>mxsalb</code> , maximum snow albedo — exactly this quantity. ERA5 does not publish it
QGRAUP	ECMWF defines 260028 <code>grle</code> and 260001 <code>tcolg</code> . ERA5 does not publish them
REFL_10CM	ECMWF defines 260134 <code>brf</code> . ERA5 does not publish it
LPI	ECMWF defines the 228050 <code>litoti</code> family, unpublished by ERA5 — and LPI is a potential index, not a flash density, so it would not be <code>litoti</code> in any case
SNOWFALLAC	it is a geometric snow depth, not a water equivalent; ECMWF's 3066 <code>sde</code> is the matching parameter and ERA5 does not publish it

The exclusions stand either way: this archive's publication rule is that a field ships only when the WRF quantity **is** the ERA5 one. But "no parameter exists" and "ERA5 does not publish it" are different claims, and only the second is true here. One field is worth naming separately. **LAI is in the model output and is not published.** CHAPTER_VARIABLES.md §6 explains at length why `lai_lv/lai_hv` are refused — the model carries one leaf area index per cell and ERA5 wants one per vegetation class, so any split would be our invention — without mentioning that the underlying field exists. It does. Note that `rdlai2d = .false.`, so it is a table lookup on the dominant category rather than an observed LAI, which is a second reason not to split it. Measured in #35.

Where the measurements live

Every number in this document was produced by a script in `tools/audit/` of the CHAPTER repository and is written up, with its method, in `docs/audit/`:

document	covers
<code>docs/audit/sample.md</code>	the 34 timesteps and why they were chosen
<code>docs/audit/harness.md</code>	how a field is compared against ERA5
<code>docs/audit/f1-pressure-levels.md</code>	§3, §4
<code>docs/audit/f2-near-surface.md</code>	§3.3
<code>docs/audit/f3-column-integrals.md</code>	§1.2, §6.2
<code>docs/audit/f4-radiation.md</code>	§1.1
<code>docs/audit/f5-water-cycle.md</code>	§5, §6.1, §8.1
<code>docs/audit/f6-soil-and-snow.md</code>	§7.1
<code>docs/audit/f7-static-and-surface.md</code>	§7.1, §8.2, §8.3
<code>docs/audit/f8-surface-exchange.md</code>	§1.3, §2, §7.2
<code>docs/audit/re-verification.md</code>	the gate the repaired converter passed before this archive was produced

Document written 2026-09-28, from measurements made 2026-09-18 to 2026-09-25.